

SAC is a off-policy algorithm. SAC trains the agent by adding entropy of the action into reward. The optimization problem becomes

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot|s_t))) \right]$$

The value function of the states becomes

$$V^{\pi}(s) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (R(s_t, a_t, s_{t+1}) + \alpha H(\pi(s))) \right]$$

The state action value is changed to include the entropy at each time step except the first one.

$$Q^{\pi}(s, a) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (R(s_t, a_t, s_{t+1})) + \sum_{t=1}^{\infty} \alpha \gamma^t H(\pi(s_t)) \mid s_0 = s, a_0 = a \right]$$

With these definition of V and Q , we have

$$V^{\pi}(s) = \mathbb{E}_{s' \in P} [R(s, a, s') + \gamma V^{\pi}(s')]$$

SAC learns a policy π_{θ} and two Q -functions Q_{ϕ_1} and Q_{ϕ_2} .
Similarities to DDPG and TD3

- both Q -functions are learned with MSE minimization, by regressing to a single shared target.
- the shared target is computed using target Q -networks and polyak interpolation,
- **clipped double- Q** trick is used to avoid overconfidence

Differences to DDPG and TD3

- target includes term from SAC's entropy regularization
- next state actions used in target come from **current** policy
- no explicit policy smooth (in DDPG and TD3 random noise is added to action for exploration). SAC trains a stochastic policy.

The entropy term and next state action term in Q^{π} are computed using the **current policy**. The next state s' is sampled from the replay buffer. But this means additional computed is need to sample the next action.

The estimate of Q -value (based one sample from experience replay) is

$$Q^{\pi} = r + \gamma(Q^{\pi}(s', \tilde{a}') - \alpha \log \pi(\tilde{a}'|s')), \tilde{a}' \sim \pi(\cdot|s')$$

We use \tilde{a}' to emphasize that next state action has to come from the current policy.

Update the Q-function

SAC trains 2 Q -function targets, the smaller one is used to compute the sample backup. (like in TD3)

The loss for the Q -networks in SAC are:

$$L(\phi_i, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',d) \sim \mathcal{D}} [(Q_{\phi_i}(s, a) - y(r, s', d))^2]$$

where

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{j=1,2} Q_{\phi_{\text{target},j}}(s', \tilde{a}') - \alpha \log \pi_{\theta}(\tilde{a}'|s') \right)$$

Update the policy

The policy should, in each state, act to maximize the expected future return plus the expected entropy, i.e. it should maximize $V^{\pi}(s)$

$$\begin{aligned} V^{\pi}(s) &= \mathbb{E}_{a \sim \pi} [Q^{\pi}(s, a)] + \alpha H(\pi(\cdot|s)) \\ &= \mathbb{E}_{a \sim \pi} [Q^{\pi}(s, a) - \alpha \log \pi(a|s)] \end{aligned}$$

Use **reparametrization trick** for policy optimization, in which a sample from $\pi(\cdot|s)$ is drawn by computing a deterministic function of state, policy params, and independent noise. We squash the Gaussian policy through tanh fun

$$\tilde{a}_{\theta}(s, \zeta) = \tanh(\mu_{\theta}(s) + \sigma_{\theta}(s) \odot \zeta), \zeta \sim N(0, 1)$$

1. tanh changes the distribution of policy (no longer Gaussian), but one can still compute the log probability. If we don't squash the Gaussian dist, then it could have infinite entropy.

2. The standard deviation depends on state. In TRPO, PPO standard deviation can be independent from the state (verify it. In my implementation, std depends on input state). According to openai, SAC does not work if the std is independent from the input state. If entropy is part of learning objective, then naturally std should be a function of input state, because entropy is a function of std.

[experiment with constant std](#)

Why reparametrization is useful? It let us write expectation over ζ rather than action, which depends on the policy parameters.

$$\mathbb{E}_{a \sim \pi_{\theta}} [Q^{\pi_{\theta}}(s, a) - \alpha \log \pi_{\theta}(a|s)] = \mathbb{E}_{\zeta \sim N} [Q^{\pi_{\theta}}(s, \tilde{a}(s, \zeta)) - \alpha \log \pi_{\theta}(\tilde{a}(s, \zeta)|s)]$$

final step to get the policy loss is to replace $Q^{\pi_{\theta}}$ with one of our function approximators Q_{ϕ_i} (the less confident one). The optimization objective becomes

$$\max_{\theta} \mathbb{E}_{s \sim \mathcal{D}, \zeta \sim N} \left[\min_{j=1,2} Q_{\phi_j}(s, \tilde{a}_{\theta}(s, \zeta)) - \alpha \log \pi_{\theta}(\tilde{a}_{\theta}(s, \zeta)|s) \right]$$

it is almost similar to DDPG's training objective.

Questions:

Why in target loss, the next state action needs to be sampled from the current policy? The intuition is with better policy, can you make a better estimate of state action value. But it behaves like a on-policy algorithm.

1 relation to generalized policy iteration

SAC can be viewed as form of GPI where both policy evaluation and policy improvement stages are softened. I think **soft** in this context just means adding entropy to make the policy more exploratory.

Soft policy evaluation can be achieved by repeated application of modified Bellman operator \mathcal{T}^π

$$\mathcal{T}^\pi := r(s_t, a_t) + \gamma \mathbb{E}_{s_{t+1} \sim p} [V_\pi(s_{t+1})]$$

Soft Bellman backup has the usual convergence property as the usual Bellman backup

Lemma 1 Define $Q^{k+1} = \mathcal{T}^\pi Q^k$, then the sequence converge to the soft Q -value of π as $k \rightarrow \infty$

$$V_\pi(s_t) = \mathbb{E}_{a_t \sim \pi} [Q(s_t, a_t) - \log \pi(a_t | s_t)]$$

In policy improvement step, we
Information projection to Gaussian family

$$\pi_{new} = \operatorname{argmin}_{\pi'} D_{KL}(\pi'(\cdot | s_t) || \frac{\exp Q^{\pi_{old}}(s_t, \cdot)}{Z^{\pi_{old}}(s_t)})$$

$$Z^{\pi_{old}}(s_t) = \int_a \exp Q^{\pi_{old}}(s_t, a) da$$

Lemma 2 Soft Policy Improvement Let $\pi_{old} \in \Pi$ and let π_{new} be the optimizer for soft policy improvement step. Then

$$Q^{\pi_{new}}(s_t, a_t) \geq Q^{\pi_{old}}(s_t, a_t)$$

for all $(s_t, a_t) \in S \times A$ with $|A| < \infty$

The proof follows the exact same pattern as in the proof of Policy Improvement Theorem.

Algorithm 1: Soft Actor-Critic

Input: initial policy params θ , Q -function params ϕ_1 and ϕ_2 , empty replay buffer \mathcal{D} .

Set target parameters equal to main parameters $\phi_{target,1} \leftarrow \phi_1, \phi_{target,2} \leftarrow \phi_2$

repeat

observe state s and select action $a \sim \pi_\theta(\cdot | s)$

execute a in the env

observe next state s', r, d signal from env

store (s, a, r, s', d) in the buffer

if s' is terminal, then reset env

if it's time to update **then**

for j in range(num of updates) **do**

randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from \mathcal{D}
 compute target for Q functions

$$y(r, s', d) = r + \gamma(1-d)(\min_{i=1,2} Q_{\phi_{target,i}}(s', \tilde{a}') - \alpha \log \pi_{\theta}(\tilde{a}'|s')), \tilde{a}' \sim \pi_{\theta}(\cdot|s')$$

update Q -fn by one step of gradient descent

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2$$

Update policy by one step gradient ascent

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} (\min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_{\theta}(s)) - \alpha \log \pi_{\theta}(\tilde{a}_{\theta}(s)|s))$$

where $\tilde{a}_{\theta}(s)$ is a sample from $\pi_{\theta}(\cdot|s)$ which is differentiable wrt θ via
 reparam trick.

update target networks

$$\phi_{target,i} \leftarrow \rho \phi_{target,i} + (1 - \rho) \phi_i$$

end for
end if
until convergence